

In [1]:

```
import pandas as pd
import datetime as dt
import mplfinance as mpf
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
import pandas_datareader.data as web
from matplotlib import style
from mplfinance.original_flavor import candlestick_ohlc
```

In [2]:

```
style.use("ggplot")
```

In [3]:

```
start = dt.datetime(2010, 2, 15) #Upper bound limit
end = dt.datetime(2020, 9, 27) #Lower bound limit
```

In [4]:

```
df = web.DataReader("AMZN", "yahoo", start, end) # Using yahoo API to pull data from S&P 500 logs
```

In [5]:

```
print(df.head(9)) #Test print the 9 most ancient entries
```

	High	Low	Open	Close	Volume	\
Date						
2010-02-16	120.500000	117.180000	120.059998	117.529999	8932700	
2010-02-17	117.129997	115.550003	117.070000	116.309998	8944800	
2010-02-18	118.510002	114.820000	115.839996	118.080002	9800100	
2010-02-19	119.089996	117.000000	117.910004	117.519997	7115600	
2010-02-22	118.970001	116.180000	117.370003	118.010002	6807300	
2010-02-23	119.250000	116.510002	118.010002	117.239998	7068200	
2010-02-24	119.800003	117.150002	117.959999	119.720001	7389900	
2010-02-25	118.339996	115.849998	118.169998	118.199997	9533400	
2010-02-26	119.430000	117.000000	117.879997	118.400002	5721600	

	Adj Close
Date	
2010-02-16	117.529999
2010-02-17	116.309998
2010-02-18	118.080002
2010-02-19	117.519997
2010-02-22	118.010002
2010-02-23	117.239998
2010-02-24	119.720001
2010-02-25	118.199997
2010-02-26	118.400002

In [6]:

```
print(df.tail(18)) #Test print the 18 most recent entries
```

	High	Low	Open	Close	Volume	\
Date						
2020-09-01	3513.870117	3467.000000	3489.580078	3499.120117	3476400	
2020-09-02	3552.250000	3486.689941	3547.000000	3531.449951	3931500	
2020-09-03	3488.409912	3303.000000	3485.000000	3368.000000	8161100	
2020-09-04	3381.500000	3111.129883	3318.000000	3294.620117	8781800	
2020-09-08	3250.850098	3130.000000	3144.000000	3149.840088	6094200	
2020-09-09	3303.179932	3185.000000	3202.989990	3268.610107	5188700	
2020-09-10	3349.889893	3170.550049	3307.219971	3175.110107	5330700	
2020-09-11	3217.340088	3083.979980	3208.689941	3116.219971	5094000	
2020-09-14	3187.389893	3096.000000	3172.939941	3102.969971	4529600	
2020-09-15	3175.020020	3108.919922	3136.159912	3156.129883	4021500	

2020-09-16	3187.239990	3074.149902	3179.989990	3078.100098	4512200
2020-09-17	3029.429932	2972.550049	3009.250000	3008.729980	6449100
2020-09-18	3037.800049	2905.540039	3031.739990	2954.909912	8892600
2020-09-21	2962.000000	2871.000000	2906.500000	2960.469971	6117900
2020-09-22	3133.989990	3000.199951	3033.840088	3128.989990	6948800
2020-09-23	3127.000000	2992.379883	3120.429932	2999.860107	5652700
2020-09-24	3069.300049	2965.000000	2977.790039	3019.790039	5529400
2020-09-25	3101.540039	2999.000000	3054.860107	3095.129883	4615200

	Adj Close
Date	
2020-09-01	3499.120117
2020-09-02	3531.449951
2020-09-03	3368.000000
2020-09-04	3294.620117
2020-09-08	3149.840088
2020-09-09	3268.610107
2020-09-10	3175.110107
2020-09-11	3116.219971
2020-09-14	3102.969971
2020-09-15	3156.129883
2020-09-16	3078.100098
2020-09-17	3008.729980
2020-09-18	2954.909912
2020-09-21	2960.469971
2020-09-22	3128.989990
2020-09-23	2999.860107
2020-09-24	3019.790039
2020-09-25	3095.129883

In [7]:

```
df.to_csv("amazon.csv") #Converts all the data into comma separated values
```

In [8]:

```
df = pd.read_csv("amazon.csv", parse_dates = True, index_col = 0) #Reads the dataframe previously saved as csv, parse the dates, and assign them as the index column {as opposed to the default [0::] indexing
```

In [9]:

```
print(df.head()) #Test print
```

	High	Low	Open	Close	Volume \
Date					
2010-02-16	120.500000	117.180000	120.059998	117.529999	8932700
2010-02-17	117.129997	115.550003	117.070000	116.309998	8944800
2010-02-18	118.510002	114.820000	115.839996	118.080002	9800100
2010-02-19	119.089996	117.000000	117.910004	117.519997	7115600
2010-02-22	118.970001	116.180000	117.370003	118.010002	6807300

	Adj Close
Date	
2010-02-16	117.529999
2010-02-17	116.309998
2010-02-18	118.080002
2010-02-19	117.519997
2010-02-22	118.010002

In [10]:

```
df.plot()
```

Out[10]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x10ee99eb8>
```

In [11]:


```
print(df["Volume"].head()) #Testing a plot call for simply the Volume column and only asking for t
```

```
the oldest five data points
```

```
Date
2010-02-16    8932700
2010-02-17    8944800
2010-02-18    9800100
2010-02-19    7115600
2010-02-22    6807300
Name: Volume, dtype: int64
```

```
In [12]:
```

```
plt.show()
```



```
In [13]:
```

```
df["90ma"] = df["Adj Close"].rolling(window=90).mean() #Calculates the moving average over 90 days
```

```
In [14]:
```

```
print(df.head()) #Test print which should NaN at the column "90ma" for all of the 5 rows printed {
it should do the same for the 90 first rows}
```

Date	High	Low	Open	Close	Volume	\
2010-02-16	120.500000	117.180000	120.059998	117.529999	8932700	
2010-02-17	117.129997	115.550003	117.070000	116.309998	8944800	
2010-02-18	118.510002	114.820000	115.839996	118.080002	9800100	
2010-02-19	119.089996	117.000000	117.910004	117.519997	7115600	
2010-02-22	118.970001	116.180000	117.370003	118.010002	6807300	

Date	Adj Close	90ma
2010-02-16	117.529999	NaN
2010-02-17	116.309998	NaN
2010-02-18	118.080002	NaN
2010-02-19	117.519997	NaN
2010-02-22	118.010002	NaN

```
In [15]:
```

```
df.dropna(inplace=True) #Will simply take out of the data frame call, the days where "90ma" is NaN
```

""To bypass the .dropna funtion, we can simply set the minimum period required to "calculate" "90ma" to zero. This allows for the column "90ma" to be populated by incremented averages of "Adj Close" from the very first day onwards. {as it reaches 90 days, the intended window calculations will take effect}

The code will look like this : `df["90ma"] = df["Adj Close"].rolling(window = 90, min_periods = 0).mean()` #Feel free to copy and test it

```
In [16]:
```

```
print(df.tail()) #Test Print
```

	High	Low	Open	Close	Volume	\
Date						
2020-09-21	2962.000000	2871.000000	2906.500000	2960.469971	6117900	
2020-09-22	3133.989990	3000.199951	3033.840088	3128.989990	6948800	
2020-09-23	3127.000000	2992.379883	3120.429932	2999.860107	5652700	
2020-09-24	3069.300049	2965.000000	2977.790039	3019.790039	5529400	
2020-09-25	3101.540039	2999.000000	3054.860107	3095.129883	4615200	

	Adj Close	90ma
Date		
2020-09-21	2960.469971	2936.810004
2020-09-22	3128.989990	2945.033781
2020-09-23	2999.860107	2951.590226
2020-09-24	3019.790039	2958.185004
2020-09-25	3095.129883	2965.360558

''' Let us now create two subplots on a 6x1 grid

- their second variables indicate where they start
- the function sharex aligns ax2's x-axis to ax1's '''

In [17]:

```
ax1 = plt.subplot2grid((6, 1), (0, 0), rowspan = 5, colspan = 1)
ax2 = plt.subplot2grid((6, 1), (5, 0), rowspan = 1, colspan = 1, sharex = ax1)
```

In [18]:


```
ax1.plot(df.index, df["Adj Close"]) #Adj Close on the first axis
ax1.plot(df.index, df["90ma"]) #90ma on the first axis
ax2.bar(df.index, df["Volume"]) #Volume on the second axis
```

Out[18]:

<BarContainer object of 2584 artists>

In [19]:

```
plt.show()
```



''' The coming part brings us all the way back to when we read and parsed the "amazon.csv" file. We are going to explore resampling {depending of the time span of our liking as to adjust data granularity} using candlestick_ohlc from the python module matplotlib.finance {we are going to set up the import at the top of the script}. We'll use the module on "Adj Close". We'll resample to every 18 days of data.

It is worth noting that appropriate OHLC data is necessary. Splits in company history, for example, can render such an analysis bumby. '''

In [20]:

```
df_ohlc = df["Adj Close"].resample("18D").ohlc() #Defining OHLC data with resampled data points for every 45 minutes
df_volume = df["Volume"].resample("18D").sum() #Resampling the volume just the same but asking for
```

```
df_volume = df_volume.groupby('Date').sum() #Resampling the volume just the same as doing for a ".sum()" as to reflect the actual {true} volume
```

In [21]:

```
print(df_ohlc.head()) #Test print of data points
```

	open	high	low	close
Date				
2010-06-23	121.449997	121.449997	108.610001	117.260002
2010-07-11	119.510002	123.650002	117.129997	117.129997
2010-07-29	116.860001	130.000000	116.860001	124.690002
2010-08-16	126.070000	135.210007	123.790001	135.210007
2010-09-03	138.789993	151.300003	137.220001	151.300003

In [22]:

```
'''
Reminder that we had already set the dates as indexes; however, with ohlc, we no longer want that.
We'll simply reset the indexes to the default [0::] format.
'''
```

Out[22]:

```
"\nReminder that we had already set the dates as indexes; however, with ohlc, we no longer want th
at. We'll simply reset the indexes to the default [0::] format.\n"
```

In [23]:

```
df_ohlc.reset_index(inplace = True)
```

In [24]:

```
print(df_ohlc.head()) #Verification
```

	Date	open	high	low	close
0	2010-06-23	121.449997	121.449997	108.610001	117.260002
1	2010-07-11	119.510002	123.650002	117.129997	117.129997
2	2010-07-29	116.860001	130.000000	116.860001	124.690002
3	2010-08-16	126.070000	135.210007	123.790001	135.210007
4	2010-09-03	138.789993	151.300003	137.220001	151.300003

In [25]:

```
df_ohlc["Date"] = df_ohlc["Date"].map(mdates.date2num) #Map out the dates after converting our curr
ent date time objects to matplotlib's mdates
```

In [26]:

```
print(df_ohlc.head()) #Test print
```

	Date	open	high	low	close
0	733946.0	121.449997	121.449997	108.610001	117.260002
1	733964.0	119.510002	123.650002	117.129997	117.129997
2	733982.0	116.860001	130.000000	116.860001	124.690002
3	734000.0	126.070000	135.210007	123.790001	135.210007
4	734018.0	138.789993	151.300003	137.220001	151.300003

In [27]:

```
ax1.xaxis_date() #Convert m_dates back to beautiful dates
```

In [28]:

```
print(df_ohlc.tail()) #Test print
```

	Date	open	high	low	close
2010-07-29	733982.0	116.860001	130.000000	116.860001	124.690002

204	737618.0	3104.000000	3196.840088	2961.969971	3033.530029
205	737636.0	3051.879883	3225.000000	3051.879883	3148.020020
206	737654.0	3182.409912	3531.449951	3182.409912	3368.000000
207	737672.0	3294.620117	3294.620117	2954.909912	2960.469971
208	737690.0	3128.989990	3128.989990	2999.860107	3095.129883

In [29]:

```
candlestick_ohlc(ax1, df_ohlc.values, width=2, colorup='g')
ax2.fill_between(df_volume.index.map(mdates.date2num), df_volume.values, 0)
```

Out[29]:

<matplotlib.collections.PolyCollection at 0x113ca6ac8>

In [30]:

```
plt.show()
```